


Notes 4: B+ Tree

Peter Bailis

Chapter 4

Indexing



Index is a **data structure** for search.

Index Types

- B-Trees (*covered next*)
 - Very good for sorted data, range search
 - *We will look at a variant called **B+ Trees***
- Hash Tables

Topics

- B-trees
- Hashing schemes


Outline:

- B-Trees \Rightarrow NEXT
- Hashing schemes


Note: This index is called B+ tree,
Sometimes just call it B-tree.

B+Tree Example


n=3



Sample non-leaf





Sample leaf node:





In textbook's notation

$n=3$

Leaf:



Non-leaf:



Size of nodes: { n+1 pointers
n keys (fixed)

Don't want nodes to be too empty

- Use at least

Non-leaf: $\lceil (n+1)/2 \rceil$ pointers

Leaf: $\lfloor (n+1)/2 \rfloor$ pointers to data





$n=3$

Non-leaf

Full node

min. node

Leaf



counts even if null

B+tree rules _____ tree of order n

- (1) All leaves at same lowest level
(balanced tree)
- (2) Pointers in leaves point to records
except for “sequence pointer”

(3) Number of pointers/keys for B+tree

| | Max ptrs | Max keys | Min ptrs→data | Min keys |
|------------------------|-------------|-------------|---------------------------|-----------------------------|
| Non-leaf (non-root) | $n+1$ | n | $\lceil (n+1)/2 \rceil$ | $\lceil (n+1)/2 \rceil - 1$ |
| Leaf (non-root) | $n+1$ | n | $\lfloor (n+1)/2 \rfloor$ | $\lfloor (n+1)/2 \rfloor$ |
| Root | $n+1$ | n | 2^* | 1 |

* When there is only one record in the B-tree, min pointers in the root is 1 (the other pointers are null)

Insert into B+tree

(a) simple case

- space available in leaf


(b) leaf overflow

(c) non-leaf overflow

(d) new root


(a) Insert key = 32

n=3




(a) Insert key = 32

n=3




(a) Insert key = 7

$n=3$




(a) Insert key = 7

n=3




(a) Insert key = 7

n=3




(c) Insert key = 160

n=3



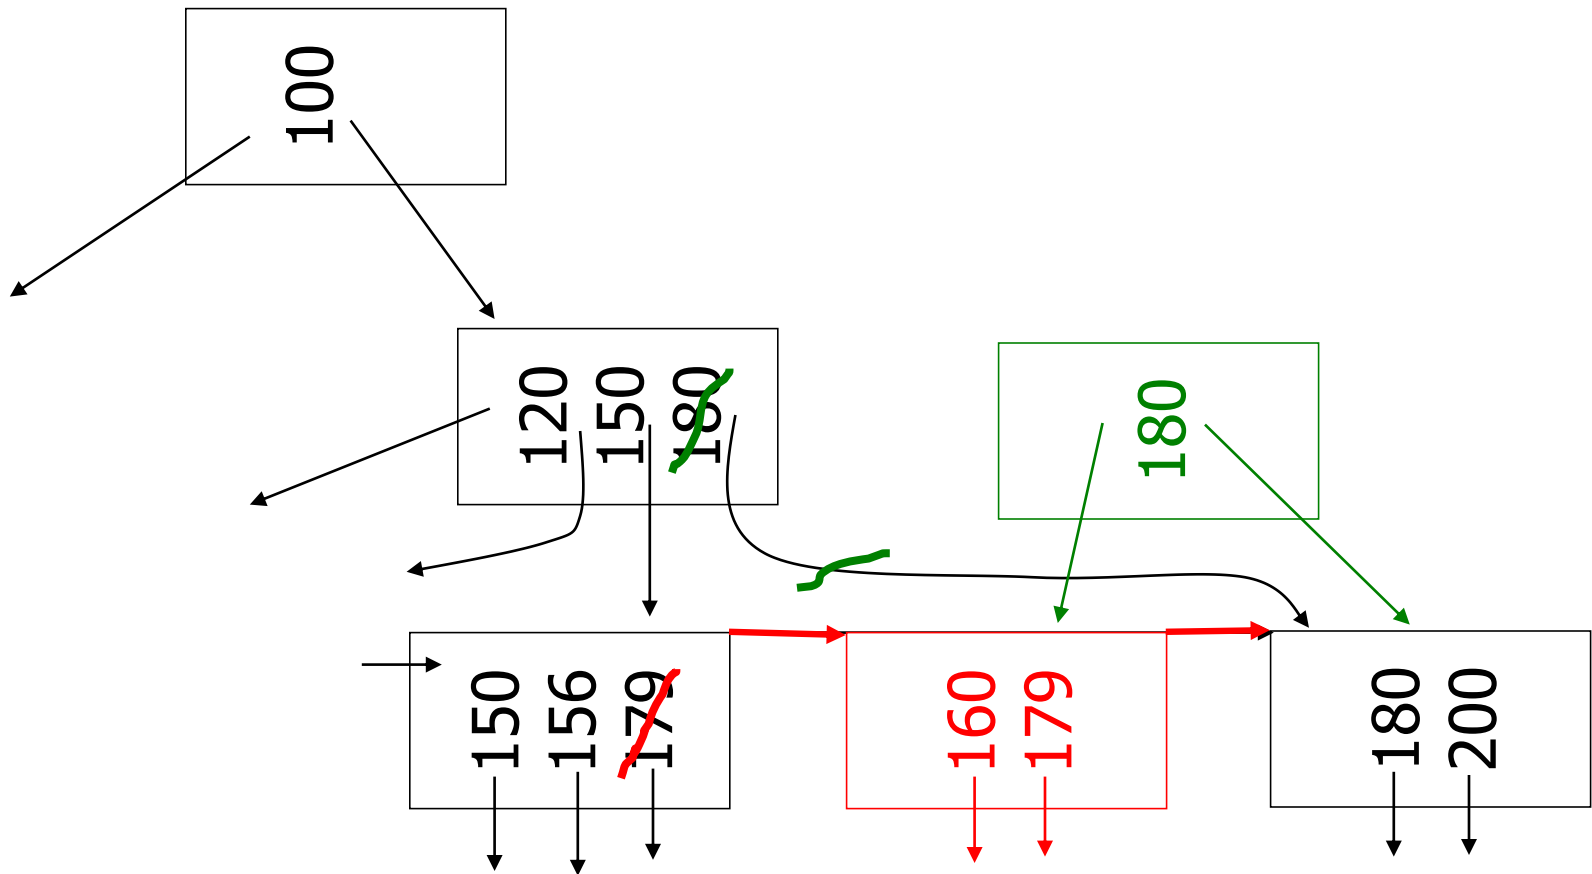
(c) Insert key = 160

n=3

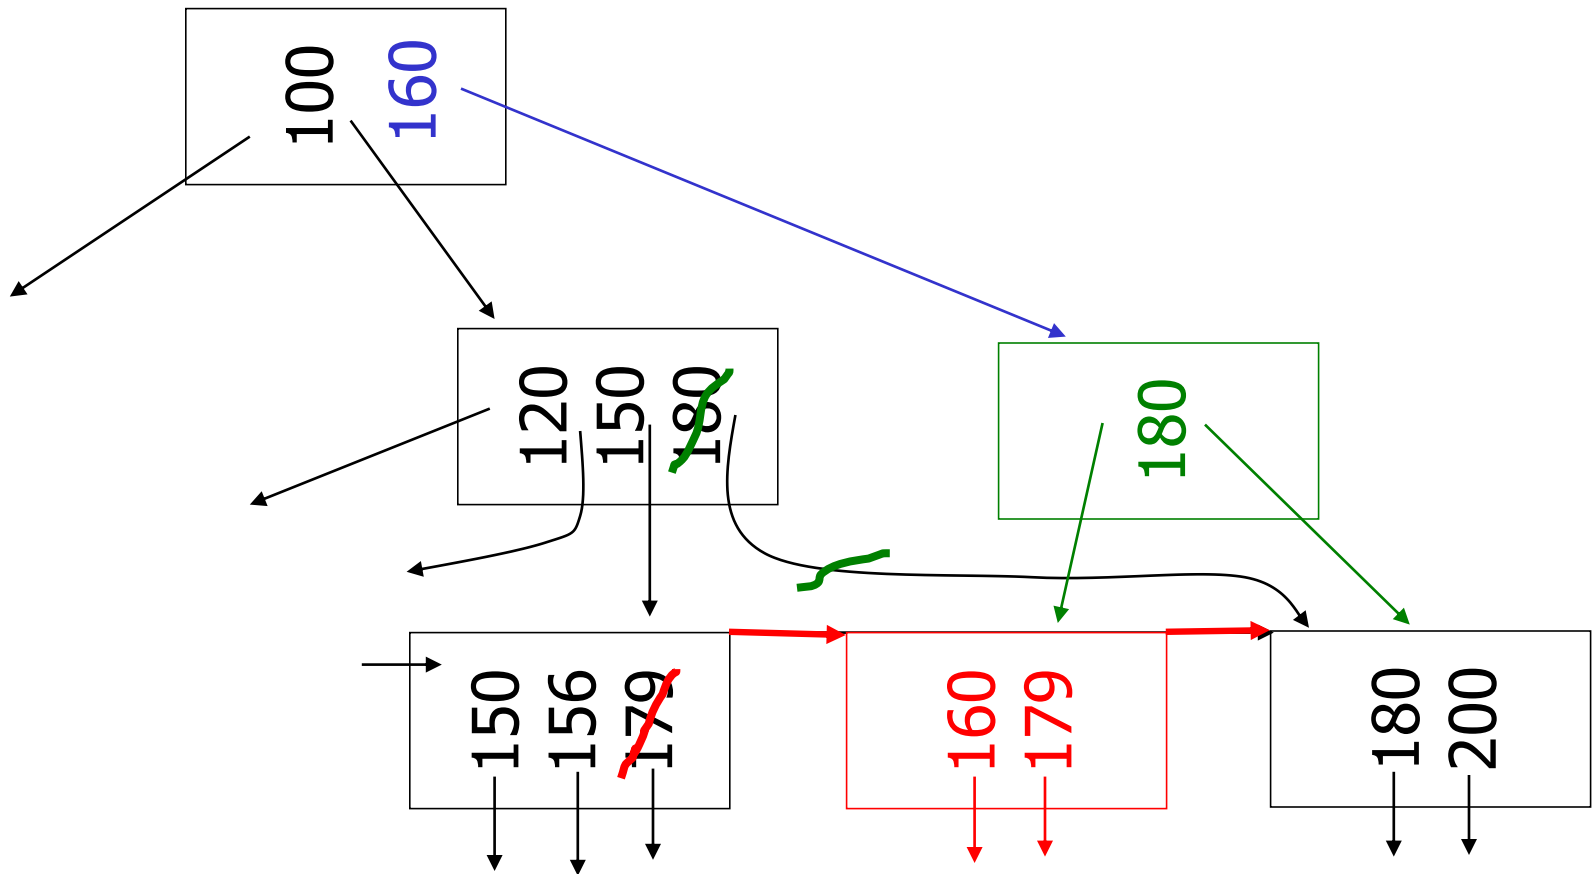


(c) Insert key = 160

n=3




n=3




(d) New root, insert 45

n=3




(d) New root, insert 45

n=3



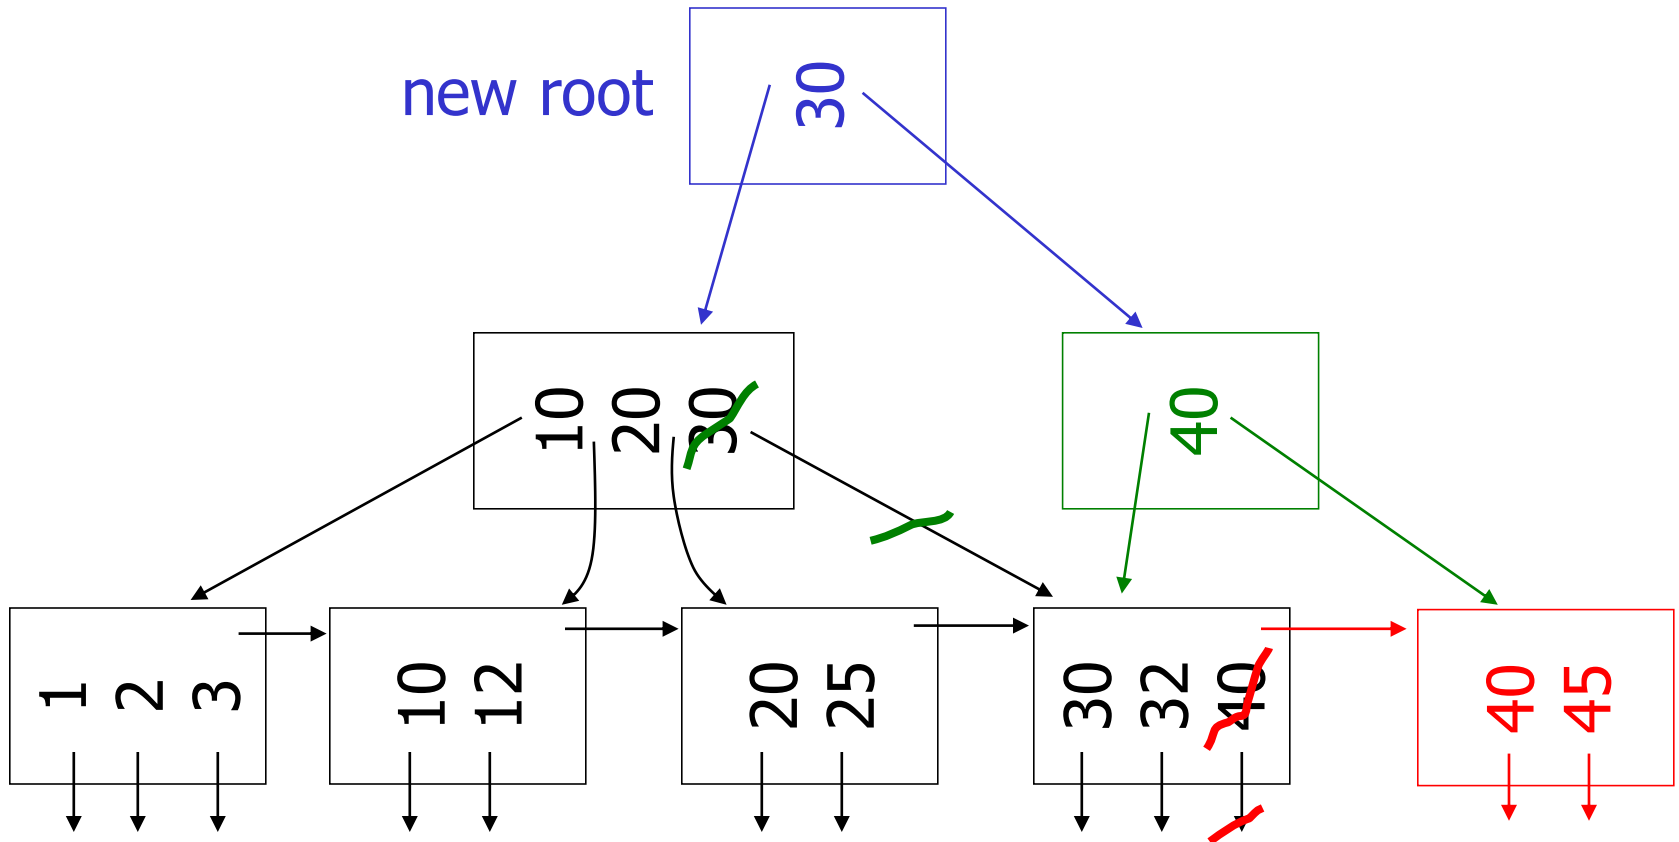
(d) New root, insert 45

n=3



(d) New root, insert 45

n=3




Deletion from B+tree

- (a) Simple case - no example
- (b) Coalesce with neighbor (sibling)
- (c) Re-distribute keys
- (d) Cases (b) or (c) at non-leaf

(b) Coalesce with sibling

– Delete 50


n=4



(b) Coalesce with sibling

– Delete 50


n=4



(c) Redistribute keys

– Delete 50


n=4



(c) Redistribute keys

– Delete 50


n=4



(d) Non-leaf coalesce

– Delete 37


n=4



(d) Non-leaf coalesce

– Delete 37


n=4



(d) Non-leaf coalesce

– Delete 37


n=4



(d) Non-leaf coalesce

– Delete 37

n=4




B+tree deletions in practice

- Often, coalescing is not implemented
 - Too hard and not worth it!

Interesting problem:

For B+tree, how large should n be?




n is number of keys / node

Variation on B+tree: B-tree (no +)

- Idea:
 - Avoid duplicate keys
 - Have record pointers in non-leaf nodes

B-tree example

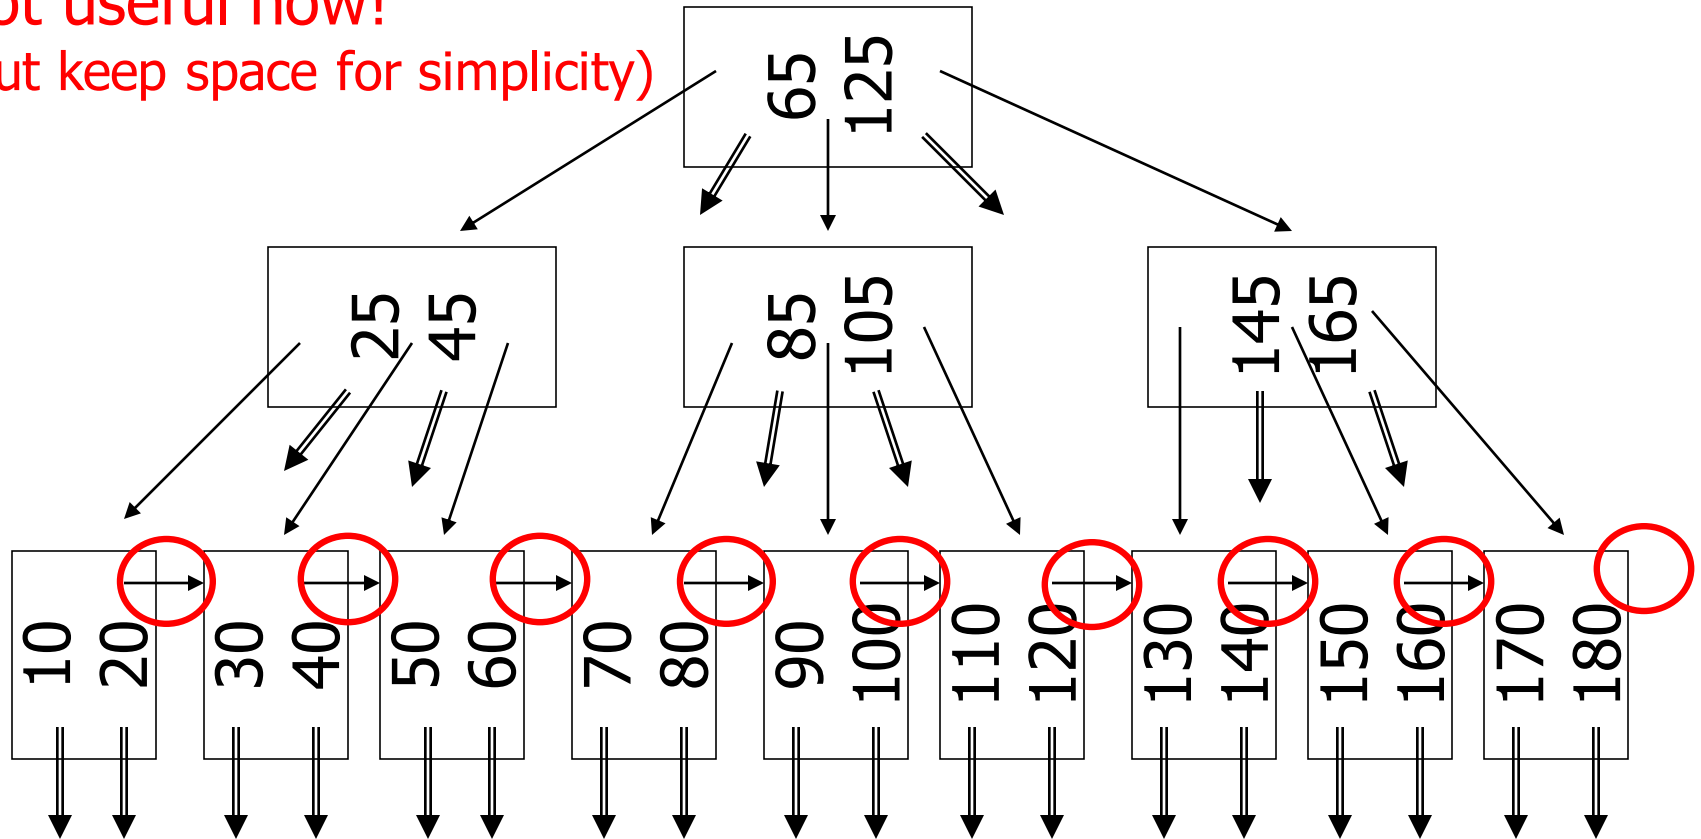
n=2



B-tree example


n=2

- sequence pointers
not useful now!
(but keep space for simplicity)



Note on inserts


- Say we insert key = 25



n=3


Note on inserts

- Say we insert key = 25



n=3

- Afterwards:



Outline/summary

- B trees
 - B+trees vs. B-trees
 - B+trees vs. indexed sequential
- Hashing schemes --> Next